

Rigorous Methodology for Finding a Validated Economic Wedge

Executive summary

The research objective should be frozen as follows:

Find an economically important, recurring job where customers already lose or spend substantial money; where software can materially change the financial outcome; where customers demonstrably pay to address the problem; where material pain remains after the best incumbent solutions are already in place; where that residual pain exists for structural rather than merely feature-level reasons; where the remaining market is large enough to support the target company; and where a new entrant can reach, onboard, and create value for customers quickly.

This is **not product-idea research** and not a search for “pain points.” The unit of research is a precisely defined **economic failure inside a workflow**:

trigger → transaction/data → decision → action → exception/failure → workaround → financial consequence → attempted resolution → final economic outcome

No candidate should be ranked until it passes seven non-compensating gates:


Hard gate	Question that must be answered	Required output
Incidence	Does the exact problem repeatedly occur in the intended ICP?	Numerator, denominator, period, confidence interval
Economic exposure	How much money is actually attributable to it?	Per-account and aggregate annual exposure
Software controllability	How much of that exposure can software actually change?	Root-cause tree + controllable fraction
Revealed willingness to pay	Do customers already sacrifice real money to address it?	Actual spend/workaround/ budget evidence
Residual grievance	Does substantial pain remain after incumbents are deployed?	Incumbent-user residual-loss evidence
Residual market capacity	Is the <i>remaining</i> monetizable market large enough?	Reproducible bottom-up ARR model

Hard gate	Question that must be answered	Required output
GTM accessibility	Can a new unknown vendor get data, prove value and get paid quickly?	Buyer/data/procurement pathway

The gates are binary. They are not a weighted score. A giant TAM cannot compensate for an unproven grievance; enthusiastic interviews cannot compensate for absent spending; incumbent complaints cannot compensate for a lack of software controllability.

For the financial objective already established in this project—₹20 crore/month in eventual profit—the business needs ₹240 crore annual profit. At mature net margins of 25%, 30%, and 40%, that corresponds to roughly **₹960 crore, ₹800 crore, and ₹600 crore of annual revenue**, respectively. A sensible default research target is therefore approximately **₹800 crore annual revenue**, while carrying ₹600–960 crore as the sensitivity range. A strong wedge/platform thesis should expose at least about **5× the required mature revenue in conservative residual ARR capacity**, implying roughly ₹4,000 crore around the ₹800 crore central case. These are decision thresholds for this search, not claims about typical SaaS companies.

The recommended funnel is:



The phrase **“up to 10” is intentional**. If only four opportunities survive all seven gates, the correct output is four—not six fabricated winners.

The evidence infrastructure needed for this is unusually good. SEC EDGAR provides free access to millions of corporate filings, more than 20 years of full-text search, and APIs covering submissions and XBRL financial data. U.S. Census County Business Patterns provides establishment counts, employment and payroll by detailed NAICS industry and size class. BLS OEWS provides occupational employment and wage estimates based on a sample of about 1.1 million establishments. O*NET contains more than 19,000 occupation-specific task statements. SAM.gov exposes federal solicitations, award notices and related procurement notices, while USAspending exposes award-level procurement data. These sources allow much more rigorous bottom-up research than generic “market size” reports. ¹

The seven hard gates

The key methodological change is to define each gate mathematically enough that two analysts should reach broadly the same conclusion from the same evidence.

Gate	Operational definition	Proposed pass criterion for a validated wedge	Automatic fail / unproven condition
Incidence	Percentage of eligible organizations experiencing the <i>exact defined failure</i> in a fixed period, plus transaction-level frequency where relevant.	Annual affected-account incidence has a 95% lower confidence bound $\geq 10\%$ in the target ICP, based on a representative sample or population-quality dataset. The exact numerator, denominator and period must be known.	Only anecdotes, reviews, vendor claims, or statistics about a broader problem; denominator unavailable.
Economic exposure	Gross annual financial consequence directly attributable to the exact failure: leakage, write-offs, excess cost, delayed cash cost, avoidable labor, penalties, downtime, lost margin, etc., with double counting removed.	P25 exposure among affected accounts must provide substantial headroom relative to the ACV required by the business model; as an initial rule, P25 gross exposure $\geq 10\times$ required ACV . Median and P25 must both be reported.	"Industry loses \$X billion" without mapping the loss to the exact root cause and ICP.
Software controllability	Fraction of economic exposure caused by root causes that software can detect, reason about, recommend on or act upon using realistically accessible data.	Root-cause model explains $\geq 80\%$ of observed exposure ; conservative controllable share is $\geq 25\%$; intervention/outcome evidence exists across ≥ 3 organizations or an equivalent high-quality study/dataset.	Majority of loss is caused by insolvency, negotiation power, physical constraints, discretionary judgment or inaccessible information that software cannot materially change.
Revealed WTP	Actual money sacrificed today to solve, contain or compensate for the problem: software, employees, BPO, auditors, consultants, fees, financing cost, penalties or outcome-based recovery fees.	Evidence of real spend from ≥ 5 named organizations, ≥ 3 contracts/awards/RFPs, or a population dataset ; expected ACV must be supported either by existing spend or a defensible fraction of measurable financial value. At least one external-spend/procurement signal is required; internal headcount alone is insufficient.	Survey respondents saying they "would pay"; vendor pricing with no evidence anybody buys; hypothetical ROI only.

Gate	Operational definition	Proposed pass criterion for a validated wedge	Automatic fail / unproven condition
Residual grievance	Financial loss/work that persists inside organizations already using the best available current solutions.	Same material residual failure found across incumbent users; qualitative saturation achieved; quantitative incumbent-user sample shows a 95% lower bound $\geq 20\%$ for affected users, or equivalent transaction evidence; at least five direct operational artifacts should confirm the residual job.	Complaint is merely missing UI/feature, or evidence comes only from non-users, old reviews or competitor marketing.
Residual market capacity	Monetizable ARR remaining <i>after</i> incidence, controllability, incumbent effectiveness, WTP and reachability are applied.	Conservative residual ARR capacity $\geq 5\times$ target annual revenue , and target revenue achievable at $\leq 15\%$ penetration of residual eligible accounts.	Market works only using headline TAM, heroic ACV, $>15\%$ penetration, or an unverified expansion assumption.
GTM accessibility	Ability of an unknown entrant to identify buyers, obtain data, deliver measurable value, get paid and expand without first replacing core infrastructure.	Buyer identifiable; historical/read-only data accessible; ≤ 10 business days from usable data to first measurable value ; target paid-entry contracting ≤ 45 days and production decision ≤ 90 days ; no core-system replacement or autonomous write access required for first value.	Pilot requires enterprise-wide deployment, regulatory license, irreversible production access, many months of integration, or buyer cannot independently authorize evaluation.

Those numerical thresholds are **research policy**, deliberately chosen to make false positives expensive. They are not universal scientific constants. The point is to stop the team moving thresholds after becoming emotionally attached to a candidate.

The crucial metric linking the gates is:

$$\text{Residual economic value/account} = E \times C \times R$$

where:

- E = gross annual economic exposure attributable to the precise problem;
- C = software-controllable share;
- R = share of that controllable exposure still unresolved after current solutions.

Then:

$$\text{Defensible ACV} = \min(\text{observed WTP ceiling, pricing benchmark, } 0.20 \times ECR)$$

The **20% value-capture ceiling** is a conservative internal rule for this search rather than an external industry fact. Using it prevents models from assuming the vendor captures most of the value created.

Finally:

$$\text{Residual ARR capacity} = N \times P \times G \times ACV$$

where:

- N = eligible organizations;
- P = lower-bound problem incidence;
- G = percentage of affected accounts that are realistically GTM-accessible;
- ACV = conservative defensible annual contract value.

This is the calculation that replaces:

“There are 50,000 companies $\times$ ₹30 lakh = ₹15,000 crore TAM.”

The methodology therefore forces the research team to answer:

How many companies have the exact problem, after competent incumbents are already installed, at a level economically large enough for us to monetize, with data we can access and a buyer we can reach?

That is the true denominator.

Evidence architecture and source hierarchy

Each factual claim should enter a central **Evidence Ledger** before it is allowed into a candidate dossier. Claims should never migrate from a source directly into a slide deck without being recorded.

A useful evidence hierarchy is:

Tier	Evidence type	How it should be used
A — direct economic evidence	Transaction logs, remittance data, claims, supplier/retailer statements, audit/recovery files, invoices, write-offs, dispute records, contract/payment data, regulator datasets	Can establish incidence, root causes, exposure and outcome directly

Tier	Evidence type	How it should be used
B — primary institutional evidence	Regulatory filings, audited reports, contracts, RFPs, government procurement awards, official job postings, customer implementation documents, vendor technical docs, earnings materials	Strong for spending, workflow architecture, buyer structure, market counts, incumbent capabilities
C — rigorous external studies	Peer-reviewed studies, government studies, association reports with disclosed samples/ methods	Good for incidence, causal/ intervention evidence and cross-company generalization
D — qualitative discovery evidence	Customer reviews, support forums, Reddit, consultant blogs, implementation communities, conference talks	Excellent for discovering hypotheses and language; weak for prevalence and TAM
E — promotional evidence	Vendor case studies, sales pages, generic analyst TAM reports	Useful for capabilities, category vocabulary and claims to verify elsewhere; never sufficient alone for a hard gate

The rule should be:

No hard gate passes from Tier D/E evidence alone.

A gate should normally have one Tier A/B source plus triangulation from at least one different source family.

Primary data sources should be prioritized aggressively. EDGAR is especially useful because researchers can search decades of corporate filings and programmatically retrieve filing/XBRL data. Census CBP can provide establishment counts by detailed industry and employment size; analysts must remember that an **establishment is not necessarily a company**, particularly for multi-unit firms. BLS OEWS can quantify the labor economics around repetitive workflows; O*NET can identify the actual tasks performed by occupations before researchers search employer job postings for evidence that those tasks persist. ²

Government procurement is particularly useful for revealed WTP. SAM.gov includes pre-solicitations, solicitations, awards and sole-source notices and exposes downloadable/API-accessible opportunity data; USAspending exposes procurement awards and attributes such as recipient, obligation, NAICS, award amount and period of performance. Those records can transform “companies seem willing to pay for this” into “organizations actually contracted for this service at these prices.” ³

Academic evidence should be used where the question is causal—particularly Gate Three. PubMed currently indexes more than 40 million biomedical citations and provides API/download tools, making it a strong source in healthcare and adjacent sciences; equivalent discipline-specific databases should be used elsewhere. ⁴

The source priority by gate should look like this:

Gate	Highest-priority sources	Secondary/discovery sources
Incidence	Transaction/claims/remittance datasets; regulator datasets; audits; peer-reviewed denominated studies	Reviews, surveys, association reports
Economic exposure	Audited financial records; transaction ledgers; filed financial statements; regulator estimates; recovery/write-off files	Industry reports, vendor case studies
Software controllability	Case-level intervention/outcome data; audits; controlled or quasi-experimental studies; workflow/event logs	Vendor performance claims, expert opinion
Revealed WTP	Signed contracts; RFPs; procurement awards; BPO SOWs; invoices; job postings + wage data; outsourcing agreements	Pricing pages, survey WTP
Residual grievance	Current incumbent-user operational data; tickets; spreadsheets; write-offs; exception queues; renewal RFPs; customer interviews	Reviews, forums, implementation consultants
Residual market capacity	Census/official registries; regulatory registries; filed company data; authoritative industry counts	Commercial databases and analyst reports
GTM accessibility	Procurement policies; previous RFP timelines; security requirements; product export/API docs; direct buyer/data-owner interviews	Vendor anecdotes and generic sales-cycle surveys

A particularly useful signal is **labor that remains despite software**. *ONET can identify the task taxonomy, company career sites show which firms are still hiring humans for those tasks, and BLS wage data can estimate the economic cost of that labor.* ONET's current database covers more than 19,000 occupation-specific task statements, while BLS OEWS publishes industry- and occupation-specific estimates from a large establishment survey. ⁵

That combination supports questions such as:

"Company has SAP + incumbent specialist X, but is still hiring 17 analysts whose descriptions say investigate exceptions, gather documents, reconcile records, contact counterparties and manually close disputes."

That is much stronger residual-grievance evidence than:

"Someone on G2 said the UI is annoying."

The **Evidence Ledger** should use this schema:

Field	Required content
Claim ID	Permanent identifier
Candidate/workflow	Exact workflow
Gate	Which hard gate it supports
Claim	One atomic falsifiable statement
ICP	Exact company/role definition
Source	Document/dataset/interview
Source tier	A-E
Date / geography	Context
Sample size	Number of accounts/transactions
Numerator	Affected units
Denominator	Eligible units
Metric	Incidence/exposure/spend/etc.
Point estimate	Numeric value
CI/range	Confidence interval or uncertainty
Extract	Exact supporting evidence
Confounders	Known limitations
Contradictory evidence	Source IDs
Analyst	Owner
Reproducibility	Query/code/file path
Decision	Supports / contradicts / neutral

Every economic number should be traceable from final table cell → model input → Evidence Ledger claim → source.

Sampling, statistics, and evidence collection

The quantitative and qualitative work should deliberately serve different purposes.

Quantitative sampling answers “how common/how large?”

Qualitative research answers “why/how exactly?”

They should never be substituted for each other.

For proportions, use **95% Wilson confidence intervals rather than naïve Wald intervals**, especially when samples are not large or incidence is away from 50%; the statistical literature has documented coverage problems with the conventional Wald interval and recommends score-based alternatives such as Wilson.

6

At the worst-case proportion of 50%, approximate sample requirements for a 95% interval are:

Independent organizations	Approx. worst-case margin of error	Intended use
43	±15 percentage points	Early directional screen
97	±10 points	Candidate validation
171	±7.5 points	Strong validation
385	±5 points	Final high-confidence incidence estimate

The important word is **organizations**.

Ten thousand transactions from one company do not constitute ten thousand independent observations about market incidence. Transaction-level datasets are excellent for understanding root causes and economic exposure, but account-level incidence should be estimated with organizations as the primary sampling unit. Where multiple transactions per firm are modeled together, use organization-level aggregation, hierarchical models or cluster-aware uncertainty estimates.

The recommended sample policy is:

Research stage	Quantitative requirement	Qualitative requirement
Broad workflow screen	Existing population data or directional evidence; no prevalence claim allowed from anecdotes	2–5 expert conversations only if required to understand workflow
Serious candidate	≥30 organizations or credible industry dataset	8–12 interviews across key roles
Seven-gate validation	≥100 independent ICP organizations, unless a population/regulatory dataset supersedes sampling	15–20 interviews per materially different homogeneous segment
Final winner	Aim for ≥385 or population-quality data on the most decision-critical incidence assumption	Continue until mechanism saturation and contradictory cases are understood

A systematic review of empirical qualitative-saturation studies found that narrowly scoped studies with relatively homogeneous populations commonly reached saturation within roughly **9–17 interviews**, while

more heterogeneous or meaning-focused studies required more. That supports a planned 15–20 interviews per homogeneous segment rather than blindly declaring “20 interviews total” adequate across several different buyer types. ⁷

For this research, use an explicit saturation rule:

Begin coding from interview one. After interview twelve, continue until **three consecutive interviews in that segment produce no new first-order root cause, workaround, incumbent limitation or buying constraint.**

The interview itself should never begin with the proposed product.

Interview block	Core question
Last occurrence	“Take me through the most recent time this exact problem happened, from trigger to final outcome.”
Frequency	“How many cases like that did you personally/your team handle last week, month and year?”
Financial outcome	“What amount was delayed, lost, overpaid, written off or spent because of that case?”
Root cause	“What specifically caused it? How was the cause proven?”
Existing stack	“Which systems, vendors and external providers touched the process?”
Residual work	“After those systems did their job, what did a human still have to do?”
Failure outcome	“What happens when the team does nothing, gets it wrong or misses the deadline?”
Spending	“How many people/vendors/tools are paid specifically for this work? Who owns that budget?”
Data	“Show me what the analyst actually opens or exports to resolve one case.”
Procurement	“Who could authorize a read-only historical analysis? What approvals would it require?”
Counterfactual	“If this job disappeared tomorrow, which cost/loss/headcount would genuinely disappear?”
Artifact request	“Can we see a redacted example: spreadsheet, claim, statement, case record, SOP, exception queue, contract or RFP?”

The most valuable interviews should therefore produce **artifacts**, not adjectives.

The root-cause data template should be:

Case ID	Org	Trigger	Failure	Root cause	\$ exposure	Current tool	Human action	Resolution	Outcome	Could software observe cause?	Could software act?
---------	-----	---------	---------	------------	-------------	--------------	--------------	------------	---------	-------------------------------	---------------------

The research team should then build a Pareto model:

$$\text{Root-cause coverage} = \frac{\text{economic exposure explained by classified root causes}}{\text{total sampled economic exposure}}$$

Gate Three should not pass until classified root causes explain at least **80%** of observed exposure. "Other" cannot contain half the money.

The **Residual-Grievance Matrix** is the central artifact for Gate Five:

Residual job	System of record	Specialist incumbent	BPO/ manual fallback	What incumbent detects	What it resolves	Human work remaining	Residual \$	Structural reason
Example pattern	ERP	Vendor A	Analysts	Exception	Partial	Gather cross-party evidence	\$X	Data boundary
Example pattern	Platform	Vendor B	BPO	Risk	No action	Contact counterparty	\$Y	Authority boundary

Only certain explanations qualify as **structural**:

Structural constraint	Meaning
Data boundary	Required truth lives across systems/companies the incumbent does not observe
Authority boundary	Incumbent can detect but cannot safely execute the resolving action
Workflow-position boundary	Product enters after the preventable error has already occurred
Incentive conflict	Incumbent profits from labor/volume that full automation would eliminate
Architecture boundary	Core system is designed to record transactions, not reconstruct economic truth
Long-tail heterogeneity	Rules/cases are too variable for incumbent's deterministic architecture
Organizational boundary	Several departments/counterparties own fragments; nobody owns outcome end-to-end

Structural constraint	Meaning
Trust/regulatory boundary	Existing automation cannot be delegated sufficient authority
Business-model boundary	Solving outcome end-to-end conflicts with vendor's product/services economics

"Vendor has not shipped feature X" **does not count**. Another product manager can ship a feature.

The research question is:

Why is the best current vendor structurally unable, economically disincentivized or positionally prevented from eliminating this residual job?

The bottom-up TAM workbook should contain no unsupported multipliers. Its core should be:

Eligible_accounts	= official/company-level account count
Incidence_low	= lower 95% confidence bound
Exposure_P25	= 25th percentile annual gross exposure
Controllable_low	= conservative software-controllable fraction
Residual_low	= conservative unsolved fraction after incumbents
Residual_value_account	= Exposure_P25 × Controllable_low × Residual_low
WTP_ceiling	= evidenced annual spend / budget
Value_price_ceiling	= 20% × Residual_value_account
Defensible_ACV benchmark)	= MIN(WTP_ceiling, Value_price_ceiling, pricing benchmark)
GTM_reachable_share	= evidence-based reachable percentage
Residual_ARR_capacity	= Eligible_accounts × Incidence_low × GTM_reachable_share × Defensible_ACV
Required_penetration	= Target_annual_revenue / Residual_ARR_capacity

Every workbook must show **conservative, base and upside** cases, but gate decisions should use the conservative case. No candidate passes because of the upside scenario.

Execution funnel and timeline

A twelve-week program is realistic **with parallel staffing**: roughly one research lead, four workflow/industry analysts, one quantitative/data analyst, and part-time interview/recruiting support. A smaller team can use the same methodology but should extend the schedule rather than reduce evidence standards.

Period	Research activity	Candidate count	Exit milestone
Week 1	Freeze ontology, scoring rules, source hierarchy, workbook and query templates	—	Research protocol locked before seeing winners
Weeks 1–2	Map economic workflows across ~20 sectors; follow revenue, spend, working capital, labor, compliance, risk and exception handling	150–200	Complete workflow universe
Week 3	Rapid public-evidence screen: identifiable buyer, observable financial consequence, preliminary incidence, evidence availability	~60	Candidates with plausible economic mechanism
Weeks 4–5	Gates One–Four: incidence, exposure, root causes, controllability, existing spending	~25	Four-gate evidence packets
Weeks 6–8	Deep incumbent reconstruction: systems of record, specialist tools, BPO, manual labor, customer residual jobs, structural causes	~15	Gate Five survivors
Weeks 8–10	Quantitative incumbent-user sampling, market-count models, account lists, data-access and procurement research	~10–12	Gates Six and Seven evaluated
Weeks 10–11	Contradiction search, expert review, model replication by second analyst, sensitivity analysis	≤10	Seven-gate validated set
Week 12	Final dossiers, comparison table, executive ranking, falsification memo	up to 10	Wedge-finding research package

The research must deliberately search for **disconfirming evidence** at every stage. Each candidate gets two analysts at the dossier stage:

Analyst A builds the strongest possible case.
Analyst B's explicit assignment is to kill it.

No analyst should be allowed to defend a market they personally discovered.

The first broad mapping should draw workflows from multiple economic pools rather than brainstorm “software ideas.” For example:

Economic pool	What the team maps
Revenue	leakage, denied revenue, underbilling, pricing errors, lost renewals
Spend	overpayments, incorrect charges, procurement leakage, supplier noncompliance

Economic pool	What the team maps
Working capital	delayed collections, inventory lock-up, payment timing, disputes
Labor	repetitive investigation, reconciliation, document assembly, chasing
Assets	downtime, maintenance exceptions, utilization losses
Risk/compliance	penalties, audit work, remediation, false positives
Financial settlement	claims, deductions, chargebacks, reconciliation, recoveries
Counterparty coordination	disputes where truth is fragmented between organizations

The first screening pass should favor **money-moving workflows with observable artifacts**, because they make rigorous validation easier.

Government and official datasets can substantially accelerate this work. Census CBP provides detailed industry and establishment-size counts, although its latest currently listed API reference year is 2023, so researchers should document the lag and supplement where fresher registries exist. SEC EDGAR provides real-time/latest filing access and more than two decades of searchable filings. SAM.gov exposes both current procurement notices and downloadable/API data. ⁸

The candidate comparison table should be mandatory and should look exactly like this:

Workflow	G1 Incidence	G2 Exposure	G3 Control	G4 WTP	G5 Residual	G6 Capacity	G7 GTM	Evidence confidence	Structural opening	Conse ARR c
Candidate A	P	P	P	P	P	P	P	A/B	Data boundary	
Candidate B	P	P	P	P	F	—	—	B	None proven	
Candidate C	P	P	C	P	P	C	P	B/C	Authority boundary	

Where:

- **P** = passed at predetermined threshold.
- **C** = conditional/unproven; cannot enter the final validated set.
- **F** = failed; research stops unless genuinely new evidence emerges.
- Confidence **A/B/C/D** reflects evidence quality separately from gate status.

Only **P/P/P/P/P/P** candidates may be called validated economic wedges.

After hard-gate qualification, and only then, surviving wedges can be ranked using softer criteria such as expansion surface, defensibility, gross margin potential, internationalization, urgency, AI leverage, platform potential, retention and founder fit.

GTM and minimally profitable MVP rules

GTM realism belongs inside wedge selection rather than being postponed until after product design.

The ideal first-customer architecture is:

historical/read-only data → independent analysis → economically measurable result → customer verifies result → customer pays → deeper integration later

The following GTM checklist should be completed during Gate Seven research:

GTM question	Pass target
Can we enumerate prospects?	Majority of ICP discoverable from registries/directories/company data
Is there a recognizable economic buyer?	Named role and budget owner known in $\geq 70\%$ of sampled accounts
Is the problem owned?	One primary team has measurable responsibility for outcome
Can data be obtained without core-system replacement?	Historical CSV/XLSX/SFTP/API/read-only access sufficient
Are required fields commonly available?	$\geq 80\%$ of sampled organizations can produce essential fields
Can a startup begin without production write access?	Yes
Can external actions remain human-approved initially?	Yes
Time from clean data to first financial finding	≤ 10 business days target
Paid entry target	≤ 45 days
Production/no-go decision target	≤ 90 days
Is a regulatory license required merely to analyze?	Prefer no
Does customer need seven departments for first pilot?	Prefer ≤ 3 mandatory stakeholders
Can value be verified objectively?	Cash recovered, cost avoided, hours eliminated, payment accelerated, etc.

The last requirement is particularly important:

The MVP's output must be an economic outcome or an auditable precursor to one—not engagement, dashboards, or “AI insights.”

Once a wedge passes all seven gates, the product-design sequence should be:

desired perfect economic outcome → perfect end-state product → minimum complete outcome loop → minimally profitable MVP

The MVP should obey these rules:

Rule	Requirement
Read-only first	Prefer exports/read-only APIs; write access only after trust
Complete loop	Solve one narrow job end-to-end instead of ten jobs halfway
Baseline	Quantify “before” economics before intervention
Auditable output	Every recommendation/action linked to evidence
Human approval	External/financial actions approved by customer until reliability is proven
No rip-and-replace	Sit beside system of record initially
Fast value	Target first verified economic finding in ≤ 10 business days
Charge from first pilots	Free pilots do not establish revealed WTP
Positive contribution	Pilot price exceeds incremental software, data, support and manually delivered analyst work
Shadow-price labor	Founder/analyst manual work counted at market cost instead of pretending it is free
Automate repeated work only	Manual MVP allowed only where it teaches what eventually becomes reusable software

The MVP may initially contain substantial human operations, but the economics must reveal whether those operations can disappear.

For pricing, test three logically different mechanisms after the wedge is chosen:

Pricing model	What it tests
Fixed paid diagnostic	Whether customer pays simply to reveal the economic opportunity
Subscription / fixed platform fee	Whether job supports recurring software budget
Base fee + outcome fee	Whether economic value is sufficiently objective for aligned pricing

A good pilot contract should define:

exact business unit/workflow; exact historical period; data fields required; customer delivery deadline; measurable success metric; fixed fee paid at signing or milestone; optional performance component; 30–45 day scope; explicit confidentiality/data-deletion terms; no external action without approval; no obligation to build customer-specific features; and an agreed production decision date.

Do **not** use “we will build whatever you need during the pilot.” That turns a wedge test into consulting.

The minimally profitable MVP should initially require **positive contribution margin on each paid engagement after shadow-pricing manual delivery labor**. A useful operating objective is to drive that toward $\geq 50\%$ contribution margin by roughly the fifth repeatable customer and toward software-like economics thereafter. Again, those are internal discipline thresholds, not claims about industry norms.

Deliverables, kill rules, and immediate actions

The research team's final package should contain **six artifacts**, all generated from the same Evidence Ledger.

Deliverable	Format	Required content
Executive decision memo	10–15 pages	Method, funnel, winners, rejected candidates, strongest conclusion, key uncertainties
Master workflow table	Spreadsheet/database	All 150–200 workflows and stage reached
Seven-gate comparison	Spreadsheet + executive table	Explicit P/C/F by gate for every serious candidate
Candidate dossiers	6–10 pages each	Exact grievance, cash flow, evidence, incumbent map, structural gap, market, GTM, falsification
Quantitative model	Reproducible workbook/notebook	Incidence CIs, exposure distributions, residual ARR model, sensitivity
Evidence repository	Database/data room	Every source, excerpt, interview, artifact and calculation

Each final dossier should follow the same structure:

Economic workflow: exact trigger-to-outcome map.

Exact grievance: one sentence precise enough to be falsifiable.

Who experiences it: ICP and buyer.

Incidence: numerator, denominator, period, CI.

Economics: P25, median, P75 exposure.

Root causes: Pareto distribution.

Controllability: controllable share and intervention evidence.

Current spending: labor/software/BPO/etc.

Incumbent landscape: system of record + specialist + services + internal labor.

Residual grievance: what remains after those solutions.

Structural reason: why it remains.

Residual market model: conservative/base/upside.

GTM path: account list, buyer, data, procurement, time-to-value.

Falsification memo: best arguments against the opportunity.

The methodology should enforce the following kill rules without exception:

Condition	Decision
No denominator for prevalence	Unproven; cannot pass Gate One
Economic statistic refers to a broader category than exact root cause	Do not use for Gate Two
Majority of monetary loss is outside software's control	Kill
WTP evidence consists only of surveys or hypothetical ROI	Fail Gate Four
Residual grievance comes only from people without modern incumbents	Fail Gate Five
Residual pain is only "incumbent lacks feature X"	Fail structural-differentiation test
Incumbent-user reviews are the only residual evidence	Unproven
Bottom-up model requires >15% residual-market penetration	Fail Gate Six
TAM relies mainly on analyst "market size" reports	Rebuild bottom-up
ACV requires capturing >20% of customer residual value without direct evidence	Reprice/model conservatively
First proof of value requires replacing ERP/core infrastructure	Fail preferred GTM architecture

Condition	Decision
First value realistically takes >90 days because of implementation rather than the intrinsic outcome	Strong GTM fail
Candidate needs hoped-for future adjacency to make current residual market adequate	Fail Gate Six
Two strong primary sources materially conflict	Use lower credible estimate until resolved
Candidate only works in upside scenario	Kill

The main research risks and controls are:

Risk	Why it creates false positives	Control
Vendor case-study bias	Best customers/outcomes are selected	Treat as Tier E for performance
Review-selection bias	Unhappy/happy users self-select	Use reviews for themes, never prevalence
Denominator mismatch	"X% of invoices" becomes "X% of companies"	Keep account and transaction incidence separate
Establishment/company confusion	Inflates target account counts	Deduplicate multi-site enterprises; Census explicitly distinguishes establishments from companies ⁹
Double-counted economics	Labor + loss + financing cost may overlap	Case-level economic accounting
Transaction clustering	Millions of records from few firms create fake precision	Organization as primary sampling unit
Broad-problem substitution	Evidence of "late payments" used for "document exceptions"	Every claim tagged to exact root cause
Labor ≠ vendor budget	Internal cost does not guarantee software spend	Require external-spend/procurement evidence
Stale market counts	Official datasets can lag	Record reference year and triangulate fresher registries
Incumbent feature creep	Apparent whitespace can disappear quickly	Require structural rather than feature gap
AI commoditization	"Better model" is copyable	Structural/data/workflow position required

Risk	Why it creates false positives	Control
International extrapolation	Rules and workflows differ by country	Validate each geography separately

The first parallel actions should be research infrastructure—not product building:

Priority	Immediate action	Purpose
Now	Freeze seven-gate definitions and thresholds in a version-controlled research protocol	Prevent goalpost movement
Now	Build Evidence Ledger + candidate comparison workbook + reproducible TAM model	Create one source of truth
Now	Build 20-sector economic-flow ontology and workflow-generation template	Produce 150–200 workflows systematically
Now	Establish automated/public-source pipelines for EDGAR, Census CBP, BLS/O*NET, SAM/USAspending and relevant regulator datasets	Make denominator/spend research repeatable; these official systems expose searchable or API-accessible data ¹⁰
Parallel	Recruit a standing research panel across CFO, controller, operations, procurement, compliance, revenue, claims and finance roles	Accelerate mechanism interviews
Parallel	Conduct 20–30 methodology-calibration interviews across unrelated workflows	Test interview coding and artifact requests, not validate a product
Parallel	Run a data-access audit: ask operators which historical exports can be produced without IT projects	Calibrate Gate Seven
Parallel	Interview procurement/security leaders about what a read-only paid analysis actually requires	Establish realistic entry constraints
Parallel	Build competitor-reconstruction template and test it on three known categories	Ensure incumbent research goes deeper than feature grids
Parallel	Have a second analyst reproduce three incidence/TAM calculations independently	Detect model/researcher error before scaling

The research team should explicitly **not** build mockups, choose technologies, invent product names, optimize landing pages or start designing a full product during this phase. Those activities create attachment to a solution before the economic position is proven.

The endpoint of this research is not:

“Here are ten promising SaaS ideas.”

It is a table in which every surviving row can support a statement of this form:

Among precisely defined organizations X, the exact failure Y affects at least P% annually with a 95% confidence interval. It causes approximately E of annual economic exposure per affected organization. At least C% of that exposure is attributable to software-controllable root causes. Customers currently spend W addressing it. Organizations already using incumbents A/B/C still experience residual loss R because of structural constraint S. Conservative bottom-up residual ARR capacity is M, enough to reach our target at $\leq 15\%$ penetration. The buyer is B, the required historical data is D, and measurable value can be demonstrated without replacing the core system within approximately ten business days.

Only when the team can write that paragraph with defensible evidence should we say we have found a validated economic wedge.

1 2 10 SEC.gov | Search Filings

<https://www.sec.gov/search-filings>

3 Contract Opportunities | SAM.gov

<https://sam.gov/opportunities>

4 PubMed

<https://pubmed.ncbi.nlm.nih.gov/>

5 O*NET OnLine

<https://www.onetonline.org/>

6 Interval Estimation for a Binomial Proportion

https://repository.upenn.edu/bitstreams/c3690c8b-efaf-485b-a984-39cfc22aae13/download?utm_source=chatgpt.com

7 Sample sizes for saturation in qualitative research: A systematic review of empirical tests - PubMed

<https://pubmed.ncbi.nlm.nih.gov/34785096/>

8 County Business Patterns (CBP) APIs

<https://www.census.gov/data/developers/data-sets/cbp-zbp/cbp-api.html>

9 County Business Patterns Methodology

https://www.census.gov/programs-surveys/cbp/technical-documentation/methodology.html?utm_source=chatgpt.com